Metodi Numerici dell'Informatica

Matrix meta-mechanics

Emanuele Rodolà
rodola@di.uniroma1.it

SAPIENZA
UNIVERSITÀ DI ROMA

2nd semester a.y. 2021/2022 · March 03, 2022

Matrix manipulation

The definition of matrix product gives rise to some alternative viewpoints that are often useful for practical manipulation of matrices.

These notes cover a few useful "meta-mechanics" of matrix products and other operations that are frequently encountered.

These results should be checked!

Do not trust everything blindly at the first exposure.
After the check, you can go blindfolded.

These notes may be updated as we go on with the course.

Transpose and inverse

A matrix $\mathbf{A}$ is symmetric if:

$$\mathbf{A} = \mathbf{A}^\top$$

If the matrix is a product $\mathbf{A} = \mathbf{BC}$, the transpose applies as follows:

$$(\mathbf{BC})^\top = \mathbf{C}^\top \mathbf{B}^\top$$

The same holds for the inverse:

$$(\mathbf{BC})^{-1} = \mathbf{C}^{-1} \mathbf{B}^{-1}$$

A matrix $\mathbf{A}$ is orthogonal if:

$$\mathbf{A}^{-1} = \mathbf{A}^\top,$$

Thus, $\mathbf{A}^\top \mathbf{A} = \mathbf{I}$ whenever $\mathbf{A}$ is orthogonal.

Products

Matrix-vector product:

$$\mathbf{X}\mathbf{y} = \begin{pmatrix} | & | & \\ \mathbf{x}_1 & \mathbf{x}_2 & \cdots \\ | & | & \end{pmatrix} \begin{pmatrix} y_1 \\ y_2 \\ \vdots \end{pmatrix} = y_1 \begin{pmatrix} | \\ \mathbf{x}_1 \\ | \end{pmatrix} + y_2 \begin{pmatrix} | \\ \mathbf{x}_2 \\ | \end{pmatrix} + \cdots$$

Products

Matrix-vector product:

$$\mathbf{X}\mathbf{y} = \begin{pmatrix} | & | & \\ \mathbf{x}_1 & \mathbf{x}_2 & \cdots \\ | & | & \end{pmatrix} \begin{pmatrix} y_1 \\ y_2 \\ \vdots \end{pmatrix} = y_1 \begin{pmatrix} | \\ \mathbf{x}_1 \\ | \end{pmatrix} + y_2 \begin{pmatrix} | \\ \mathbf{x}_2 \\ | \end{pmatrix} + \cdots$$

Vector-matrix product; it's just a transposed version of the above:

$$\mathbf{z}^\top \mathbf{A} = (\mathbf{A}^\top \mathbf{z})^\top$$

Products

Matrix-vector product:

$$\mathbf{X}\mathbf{y} = \begin{pmatrix} | & | & \\ \mathbf{x}_1 & \mathbf{x}_2 & \cdots \\ | & | & \end{pmatrix} \begin{pmatrix} y_1 \\ y_2 \\ \vdots \end{pmatrix} = y_1 \begin{pmatrix} | \\ \mathbf{x}_1 \\ | \end{pmatrix} + y_2 \begin{pmatrix} | \\ \mathbf{x}_2 \\ | \end{pmatrix} + \cdots$$

Vector-matrix product; it's just a transposed version of the above:

$$\mathbf{z}^\top \mathbf{A} = (\mathbf{A}^\top \mathbf{z})^\top$$

Matrix-matrix product:

$$\mathbf{X}\mathbf{Y} = \begin{pmatrix} — & \mathbf{x}_1^\top & — \\ — & \mathbf{x}_2^\top & — \\ & \vdots & \end{pmatrix} \begin{pmatrix} | & | & \\ \mathbf{y}_1 & \mathbf{y}_2 & \cdots \\ | & | & \end{pmatrix} =$$

Products

Matrix-vector product:

$$\mathbf{X}\mathbf{y} = \begin{pmatrix} | & | & \\ \mathbf{x}_1 & \mathbf{x}_2 & \cdots \\ | & | & \end{pmatrix} \begin{pmatrix} y_1 \\ y_2 \\ \vdots \end{pmatrix} = y_1 \begin{pmatrix} | \\ \mathbf{x}_1 \\ | \end{pmatrix} + y_2 \begin{pmatrix} | \\ \mathbf{x}_2 \\ | \end{pmatrix} + \cdots$$

Vector-matrix product; it's just a transposed version of the above:

$$\mathbf{z}^\top \mathbf{A} = (\mathbf{A}^\top \mathbf{z})^\top$$

Matrix-matrix product:

$$\mathbf{X}\mathbf{Y} = \begin{pmatrix} — & \mathbf{x}_1^\top & — \\ — & \mathbf{x}_2^\top & — \\ & \vdots & \end{pmatrix} \begin{pmatrix} | & | & \\ \mathbf{y}_1 & \mathbf{y}_2 & \cdots \\ | & | & \end{pmatrix} = \begin{pmatrix} | & | & \\ \mathbf{X}\mathbf{y}_1 & \mathbf{X}\mathbf{y}_2 & \cdots \\ | & | & \end{pmatrix}$$

$$= \begin{pmatrix} — & \mathbf{x}_1^\top \mathbf{Y} & — \\ — & \mathbf{x}_2^\top \mathbf{Y} & — \\ & \vdots & \end{pmatrix}$$

Products

Matrix-matrix product:

$$\mathbf{X}\mathbf{Y} = \begin{pmatrix} — & \mathbf{x}_1^\top & — \\ — & \mathbf{x}_2^\top & — \\ & \vdots & \end{pmatrix} \begin{pmatrix} | & | & \\ \mathbf{y}_1 & \mathbf{y}_2 & \cdots \\ | & | & \end{pmatrix} = \begin{pmatrix} & \cdots & \\ \vdots & \mathbf{x}_i^\top \mathbf{y}_j & \vdots \\ & \cdots & \end{pmatrix}$$

Products

Matrix-matrix product:

$$\mathbf{X}\mathbf{Y} = \begin{pmatrix} — & \mathbf{x}_1^\top & — \\ — & \mathbf{x}_2^\top & — \\ & \vdots & \end{pmatrix} \begin{pmatrix} | & | & \\ \mathbf{y}_1 & \mathbf{y}_2 & \cdots \\ | & | & \end{pmatrix} = \begin{pmatrix} & \cdots & \\ \vdots & \mathbf{x}_i^\top \mathbf{y}_j & \vdots \\ & \cdots & \end{pmatrix}$$

Vector-vector product (inner):

$$\mathbf{x}^\top \mathbf{y} = \alpha$$

Products

Matrix-matrix product:

$$\mathbf{X}\mathbf{Y} = \begin{pmatrix} — & \mathbf{x}_1^\top & — \\ — & \mathbf{x}_2^\top & — \\ & \vdots & \end{pmatrix} \begin{pmatrix} | & | & \\ \mathbf{y}_1 & \mathbf{y}_2 & \cdots \\ | & | & \end{pmatrix} = \begin{pmatrix} & \cdots & \\ \vdots & \mathbf{x}_i^\top \mathbf{y}_j & \vdots \\ & \cdots & \end{pmatrix}$$

Vector-vector product (inner):

$$\mathbf{x}^\top \mathbf{y} = \alpha$$

Vector-vector product (outer):

$$\mathbf{x}\mathbf{y}^\top = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \end{pmatrix} \begin{pmatrix} y_1 & y_2 & \cdots \end{pmatrix} = \begin{pmatrix} | & | & \\ y_1\mathbf{x} & y_2\mathbf{x} & \cdots \\ | & | & \end{pmatrix}$$

For example, a matrix full of ones is just $\mathbf{1}\mathbf{1}^\top$.

Diagonal matrices

Matrix-vector product:

$$\begin{pmatrix} d_1 & & \\ & d_2 & \\ & & \ddots \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ \vdots \end{pmatrix} = \begin{pmatrix} d_1 x_1 \\ d_2 x_2 \\ \vdots \end{pmatrix}$$

Diagonal matrices

Matrix-vector product:

$$\begin{pmatrix} d_1 & & \\ & d_2 & \\ & & \ddots \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ \vdots \end{pmatrix} = \begin{pmatrix} d_1 x_1 \\ d_2 x_2 \\ \vdots \end{pmatrix}$$

Matrix-matrix product:

$$\begin{pmatrix} d_1 & & \\ & d_2 & \\ & & \ddots \end{pmatrix} \begin{pmatrix} — & \mathbf{x}_1^\top & — \\ — & \mathbf{x}_2^\top & — \\ & \vdots & \end{pmatrix} = \begin{pmatrix} — & d_1 \mathbf{x}_1^\top & — \\ — & d_2 \mathbf{x}_2^\top & — \\ & \vdots & \end{pmatrix}$$

Diagonal matrices

Matrix-vector product:

$$\begin{pmatrix} d_1 & & \\ & d_2 & \\ & & \ddots \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ \vdots \end{pmatrix} = \begin{pmatrix} d_1 x_1 \\ d_2 x_2 \\ \vdots \end{pmatrix}$$

Matrix-matrix product:

$$\begin{pmatrix} d_1 & & \\ & d_2 & \\ & & \ddots \end{pmatrix} \begin{pmatrix} — & \mathbf{x}_1^\top & — \\ — & \mathbf{x}_2^\top & — \\ & \vdots & \end{pmatrix} = \begin{pmatrix} — & d_1\mathbf{x}_1^\top & — \\ — & d_2\mathbf{x}_2^\top & — \\ & \vdots & \end{pmatrix}$$

From the other side:

$$\begin{pmatrix} | & | & \\ \mathbf{y}_1 & \mathbf{y}_2 & \cdots \\ | & | & \end{pmatrix} \begin{pmatrix} d_1 & & \\ & d_2 & \\ & & \ddots \end{pmatrix} = \begin{pmatrix} | & | & \\ d_1\mathbf{y}_1 & d_2\mathbf{y}_2 & \cdots \\ | & | & \end{pmatrix}$$

Trace

The trace of $\mathbf{A}$ is the sum of its diagonal elements:

$$\operatorname{tr}(\mathbf{A}) = \sum_i a_{ii}$$

It is a linear mapping, since:

$$\operatorname{tr}(\mathbf{A} + \mathbf{B}) = \operatorname{tr}(\mathbf{A}) + \operatorname{tr}(\mathbf{B})$$
$$\operatorname{tr}(a\mathbf{A}) = a \operatorname{tr}(\mathbf{A})$$

Trace

The trace of $\mathbf{A}$ is the sum of its diagonal elements:

$$\operatorname{tr}(\mathbf{A}) = \sum_i a_{ii}$$

It is a linear mapping, since:

$$\operatorname{tr}(\mathbf{A} + \mathbf{B}) = \operatorname{tr}(\mathbf{A}) + \operatorname{tr}(\mathbf{B})$$
$$\operatorname{tr}(a\mathbf{A}) = a \operatorname{tr}(\mathbf{A})$$

It is invariant to the transpose:

$$\operatorname{tr}(\mathbf{A}) = \operatorname{tr}(\mathbf{A}^\top)$$

Trace

The trace of $\mathbf{A}$ is the sum of its diagonal elements:

$$\operatorname{tr}(\mathbf{A}) = \sum_i a_{ii}$$

It is a linear mapping, since:

$$\operatorname{tr}(\mathbf{A} + \mathbf{B}) = \operatorname{tr}(\mathbf{A}) + \operatorname{tr}(\mathbf{B})$$
$$\operatorname{tr}(a\mathbf{A}) = a\operatorname{tr}(\mathbf{A})$$

It is invariant to the transpose:

$$\operatorname{tr}(\mathbf{A}) = \operatorname{tr}(\mathbf{A}^\top)$$

It is invariant to cyclic permutations:

$$\operatorname{tr}(\mathbf{ABC}) = \operatorname{tr}(\mathbf{CAB}) = \operatorname{tr}(\mathbf{BCA})$$

Norms

The (squared) Frobenius norm for a matrix $\mathbf{X}$ is:

$$\|\mathbf{X}\|_F^2 = \mathrm{tr}(\mathbf{X}\mathbf{X}^\top) = \mathrm{tr}(\mathbf{X}^\top\mathbf{X})$$

Norms

The (squared) Frobenius norm for a matrix $\mathbf{X}$ is:

$$\|\mathbf{X}\|_F^2 = \mathrm{tr}(\mathbf{X}\mathbf{X}^\top) = \mathrm{tr}(\mathbf{X}^\top\mathbf{X})$$

If $\mathbf{X}$ is a vector, this reduces to:

$$\|\mathbf{x}\|_2^2 = \mathrm{tr}(\underbrace{\mathbf{x}\mathbf{x}^\top}_{n\times n}) = \mathrm{tr}(\underbrace{\mathbf{x}^\top\mathbf{x}}_{1\times 1}) = \mathbf{x}^\top\mathbf{x}$$

Norms

The (squared) Frobenius norm for a matrix $\mathbf{X}$ is:

$$\|\mathbf{X}\|_F^2 = \mathrm{tr}(\mathbf{X}\mathbf{X}^\top) = \mathrm{tr}(\mathbf{X}^\top\mathbf{X})$$

If $\mathbf{X}$ is a vector, this reduces to:

$$\|\mathbf{x}\|_2^2 = \mathrm{tr}(\underbrace{\mathbf{x}\mathbf{x}^\top}_{n\times n}) = \mathrm{tr}(\underbrace{\mathbf{x}^\top\mathbf{x}}_{1\times 1}) = \mathbf{x}^\top\mathbf{x}$$

For example, for matrices $\mathbf{A}$ and $\mathbf{B}$ we can derive the distance:

$$\begin{aligned}\|\mathbf{A}-\mathbf{B}\|_F^2 &= \mathrm{tr}((\mathbf{A}-\mathbf{B})^\top(\mathbf{A}-\mathbf{B})) \\ &= \mathrm{tr}(\mathbf{A}^\top\mathbf{A}) - 2\mathrm{tr}(\mathbf{A}^\top\mathbf{B}) + \mathrm{tr}(\mathbf{B}^\top\mathbf{B})\,,\end{aligned}$$

where we used the linearity of the trace and its invariance to transposition.

Ones

A vector $\mathbf{1}$ of ones can be used to calculate sums easily.

Sum up the elements of $\mathbf{A}$ along each row:

$$\mathbf{A}\mathbf{1}$$

Sum up along each column:

$$\mathbf{1}^\top \mathbf{A} = (\mathbf{A}^\top \mathbf{1})^\top$$

Ones

A vector $\mathbf{1}$ of ones can be used to calculate sums easily.

Sum up the elements of $\mathbf{A}$ along each row:

$$\mathbf{A}\mathbf{1}$$

Sum up along each column:

$$\mathbf{1}^\top \mathbf{A} = (\mathbf{A}^\top \mathbf{1})^\top$$

Sum up all the elements of $\mathbf{A}$:

$$\mathbf{1}^\top \mathbf{A}\mathbf{1}$$

Ones

A vector $\mathbf{1}$ of ones can be used to calculate sums easily.

Sum up the elements of $\mathbf{A}$ along each row:

$$\mathbf{A}\mathbf{1}$$

Sum up along each column:

$$\mathbf{1}^\top \mathbf{A} = (\mathbf{A}^\top \mathbf{1})^\top$$

Sum up all the elements of $\mathbf{A}$:

$$\mathbf{1}^\top \mathbf{A}\mathbf{1}$$

Note the following relationship:

$$\mathbf{x}^\top \mathbf{A}\mathbf{x} = \mathrm{tr}(\mathbf{x}^\top \mathbf{A}\mathbf{x}) = \mathrm{tr}(\mathbf{x}\mathbf{x}^\top \mathbf{A}) = \mathrm{tr}(\mathbf{X}\mathbf{A})$$

Ones

A vector $\mathbf{1}$ of ones can be used to calculate sums easily.

Sum up the elements of $\mathbf{A}$ along each row:

$$\mathbf{A1}$$

Sum up along each column:

$$\mathbf{1}^\top \mathbf{A} = (\mathbf{A}^\top \mathbf{1})^\top$$

Sum up all the elements of $\mathbf{A}$:

$$\mathbf{1}^\top \mathbf{A1}$$

Note the following relationship:

$$\mathbf{x}^\top \mathbf{Ax} = \mathrm{tr}(\mathbf{x}^\top \mathbf{Ax}) = \mathrm{tr}(\mathbf{xx}^\top \mathbf{A}) = \mathrm{tr}(\mathbf{XA})$$

Constructing the matrix $\mathbf{xx}^\top$ from the vector $\mathbf{x}$ is also called lifting.

Permutation matrices

$$\begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} 1 \\ 2 \\ 3 \end{pmatrix} = \begin{pmatrix} 2 \\ 1 \\ 3 \end{pmatrix}$$

Permutation matrices

$$\begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} 1 \\ 2 \\ 3 \end{pmatrix} = \begin{pmatrix} 2 \\ 1 \\ 3 \end{pmatrix}$$

Permutation matrices are orthogonal.

Their product is still a permutation matrix.

Permutation matrices

$$\begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} 1 \\ 2 \\ 3 \end{pmatrix} = \begin{pmatrix} 2 \\ 1 \\ 3 \end{pmatrix}$$

Permutation matrices are orthogonal.

Their product is still a permutation matrix.

Their convex combination is a doubly stochastic matrix:

$$\alpha \mathbf{P} + (1-\alpha)\mathbf{Q} = \mathbf{D} \quad \text{with } \alpha \in [0,1]$$

that is, we get $\mathbf{D1} = \mathbf{1}$ and $\mathbf{D}^\top \mathbf{1} = \mathbf{1}$.

Gradients of traces

The following expression appears frequently in practice:

$$\nabla \mathrm{tr}(\mathbf{A}) = \nabla \sum_i a_{ii}$$

which requires the computation of the partial derivatives:

$$\frac{\partial}{\partial a_{ij}} \sum_i a_{ii}$$

A common pitfall is the following invalid operation:

$$\underbrace{\nabla \mathrm{tr}(\mathbf{A})}_{\substack{\text{gradient of a} \\ \text{scalar function}}} \Rightarrow \mathrm{tr}(\underbrace{\nabla \mathbf{A}}_{\text{undefined}})$$

Also observe that $\nabla \mathrm{tr}(\mathbf{A})$ is a matrix, while $\mathrm{tr}(\cdots)$ is a scalar.

Suggested reading

For a review of matrix calculus, read Chapters 0.0 – 0.2 of the book:

R. Horn & C. Johnson, "Matrix Analysis - 2nd ed". Cambridge University Press, 2013